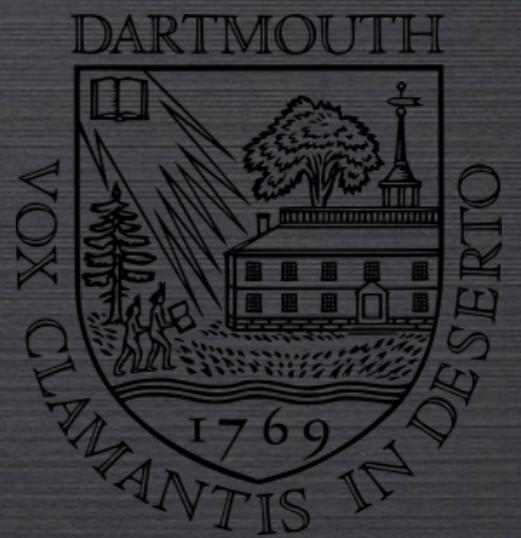


# ELF ECCENTRICITIES

JULIAN BANGERT, REBECCA SHAPIRO,  
SERGEY BRATUS

DARTMOUTH COLLEGE



# THIS TALK BROUGHT TO YOU BY THE LETTER “L”

---

- ELF: Executable Linkable Format
- Linking/loading is a complex computation on complex inputs
- Linking/loading paints magical memory trees (vma\_structs, ...)
- Parsed by exec(), ld.so, IDA, ... - each slightly differently
- Full of mystery? “Invisible”?

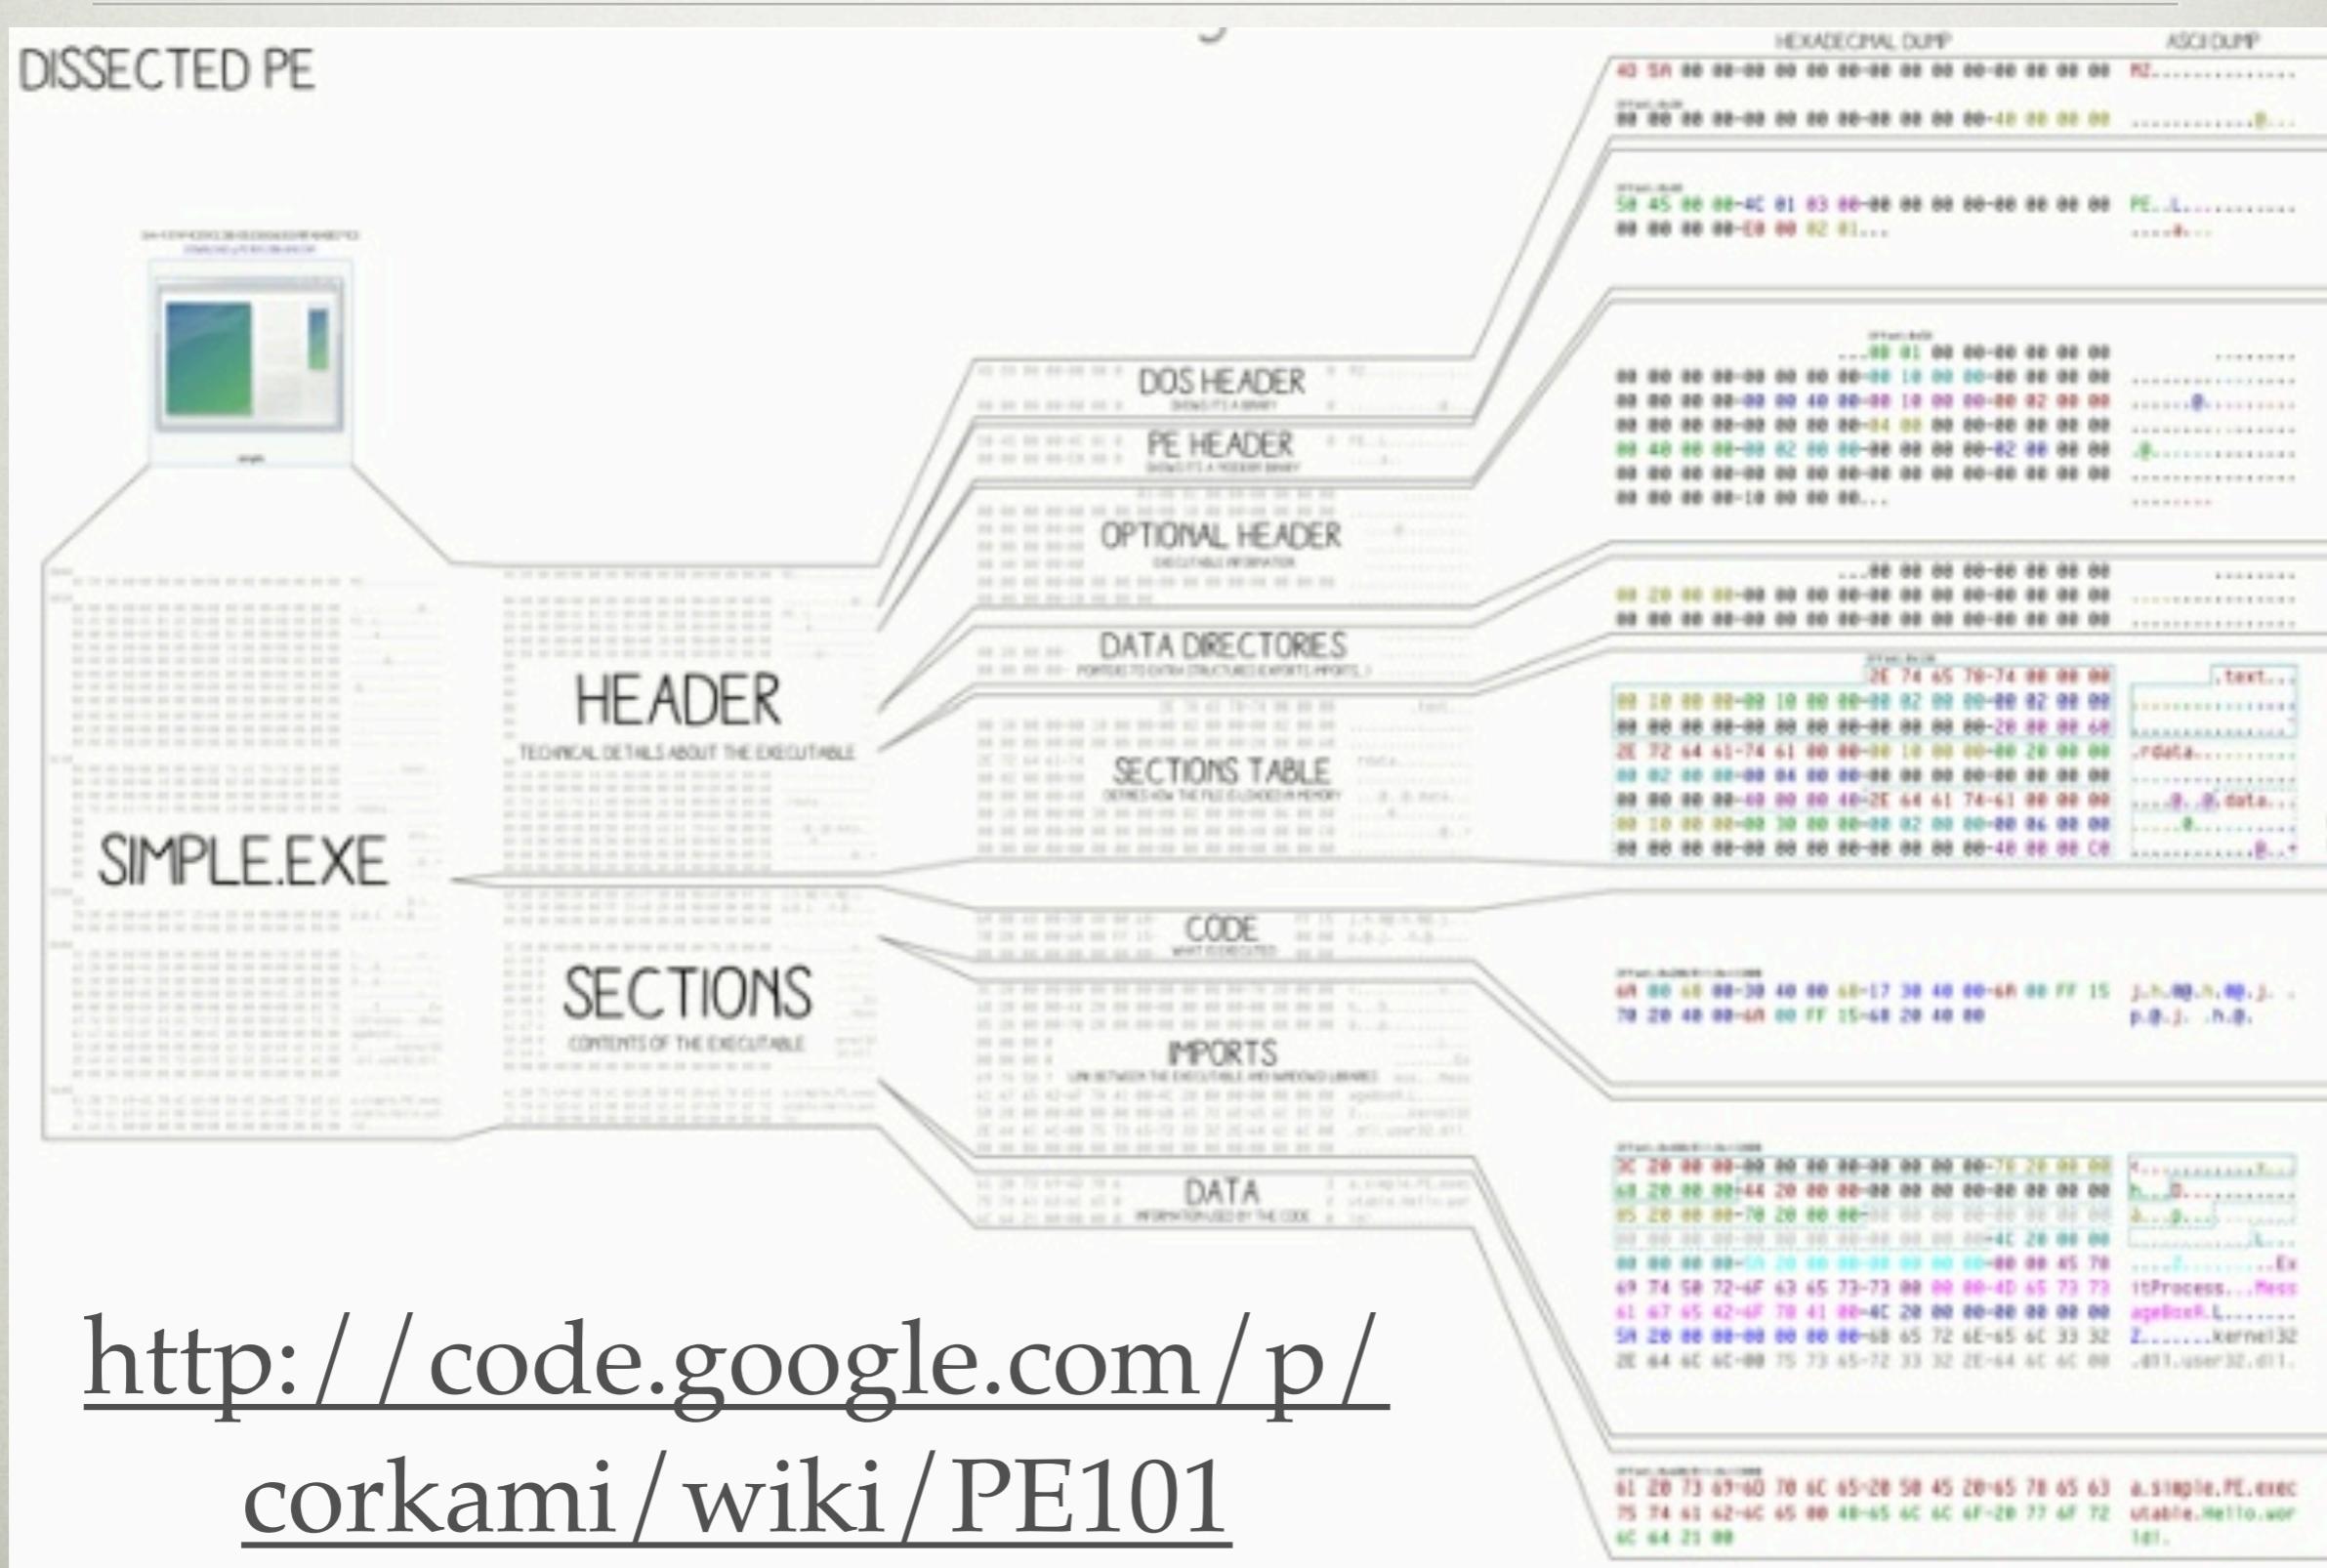


# ELFs IN ALL THE THINGS

---

- Executables, .so, .o, core dumps, ...
- Replaced non-extensible AOUT
- Designed by System V working group, extremely extensible:
  - keeps code/data **intelligible** (sections, symbols, dynamic symbols)
  - allows code/data to **compose** (“link”)
  - allows code/data to **mutate** (relocs ~ASLR)

# PE, A COUSIN OF ELF



<http://code.google.com/p/corkami/wiki/PE101>

# MACH-O, A COUSIN OF ELF

---

“This space intentionally left blank”

More to come from @bxsays

# FURTHER REFLECTIONS ON TRUSTING TRUST (I)

---

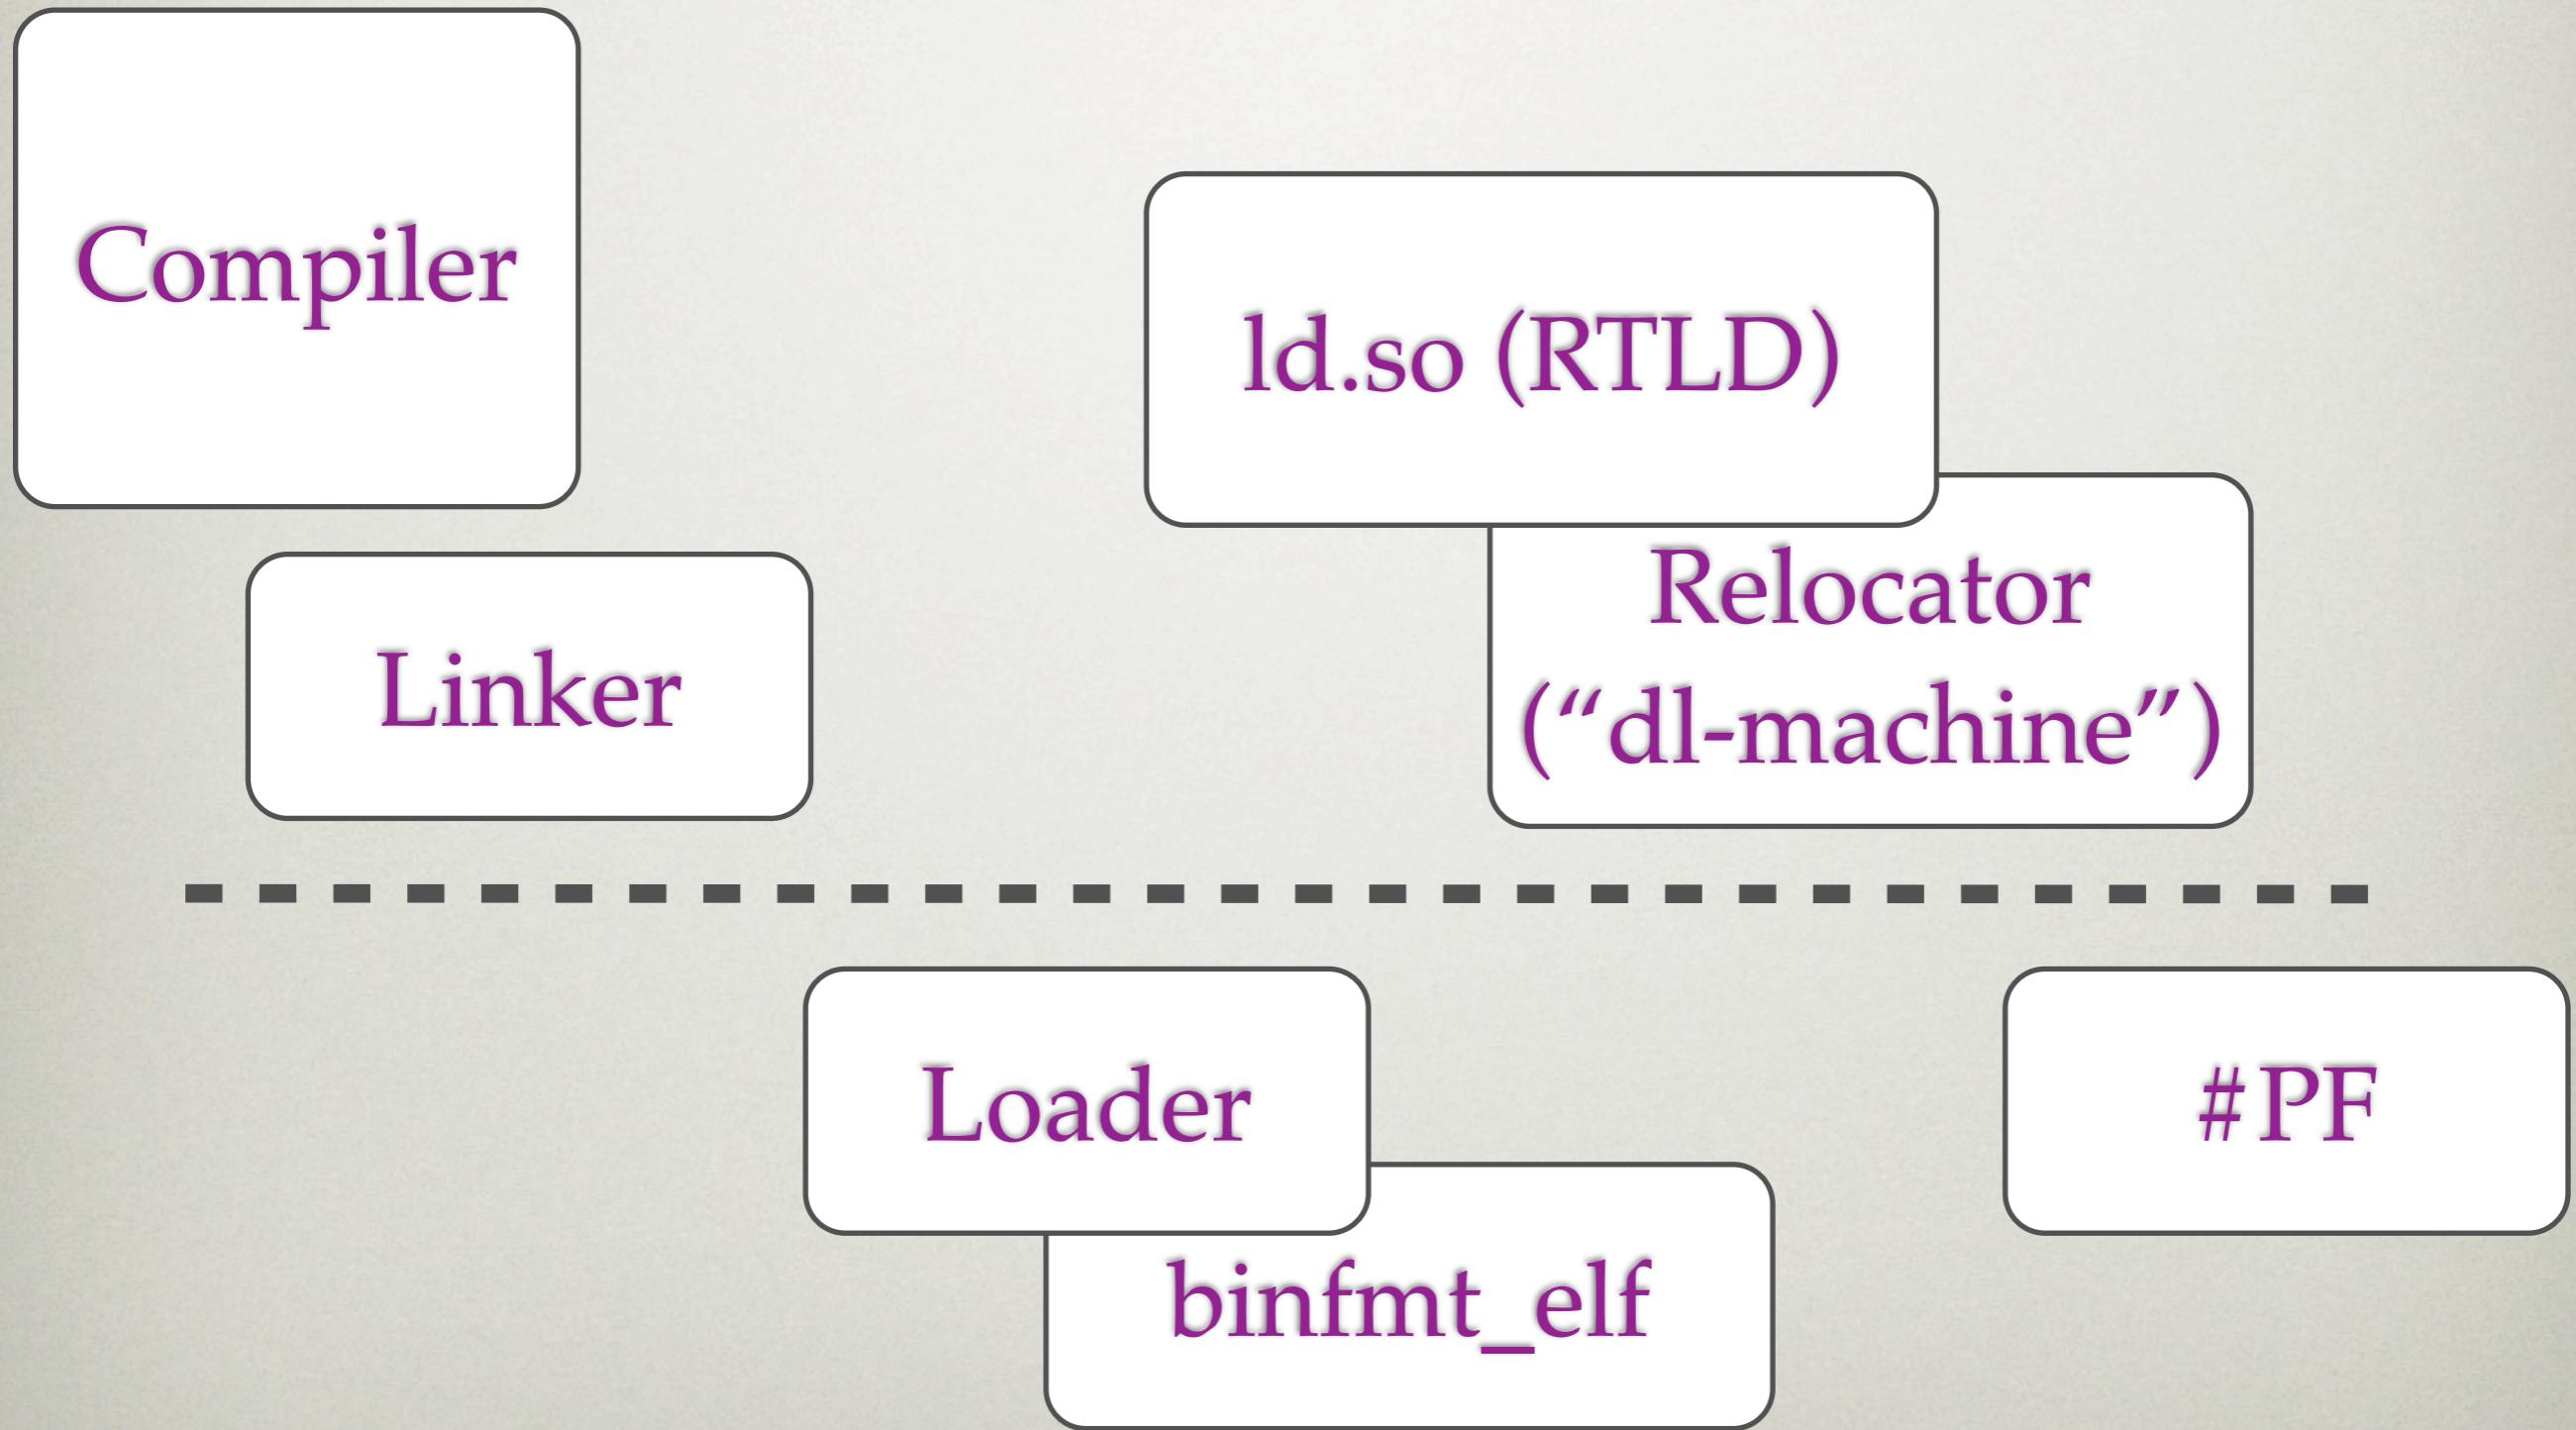
- Ken Thompson, “Reflections on Trusting Trust”, 1984
  - *“You can't trust code that you did not totally create yourself”*
  - many links in the trust chain are ‘invisible’ (down to “*well-installed microcode bugs*”)

# FURTHER REFLECTIONS ON TRUSTING TRUST (II)

---

- What if there were **no bugs** in any given piece of sw/hw link of the trust chain?
  - What could go wrong then?
- What are the **execution models** / computational strength of the links?
  - What can one do with **well-formed** crafted inputs?

# THE ELF/ABI CASE STUDY



# A WHIRLWIND OF WEIRD MACHINES

---

- Relocation entries + dynamic symbols = Turing machine on process' address space
- DWARF exception handling data + .eh\_frame + Glibc = Turing machine
- GDT + IDT + TSS + page tables + #PF + #DF = Turing machine in ia32
- More coming :)

# LANGSEC ANGLE (1)

---

- “Any Input Is A Program”
  - Input **executes** on input handlers (drives state changes & transitions)
  - Only a **well-defined** execution model can be trusted
- Input handler is either a **recognizer** for the inputs as a well-defined language, or it’s pwned (see <http://langsec.org/>)

# LANGSEC ANGLE (2)

---

- Parse differentials create subtle vulns
- One input, two parsers: A Tale of Woe
  - “*PKI Layer Cake*” X.509 parser differentials,  
Dan Kaminsky, Meredith L. Patterson, Len Sassaman
  - Cert parser differentials between CAs and clients  
broke the X.509 trust model as we know it

# LANGSEC ANGLE (3)

---

- Two views of the same data => confusion
  - Redundant data: which version is taken?
  - How many ELF parsers in ABI toolchain?  
Kernel's `binfmt_elf` loader vs user `ld.so`

DER TRAGÖDIE ERSTER TEIL

TOUR DE ELF  
(WITH A DASH OF FAIL)

# A BEAUTIFUL DREAM...

---

- Tools don't need different parsers for libraries, object files and executables
- Tools work for all platforms
- Postel's principle: Allow tools to ignore information they don't understand, change tools only as needed
- Otherwise, no .so or ASLR

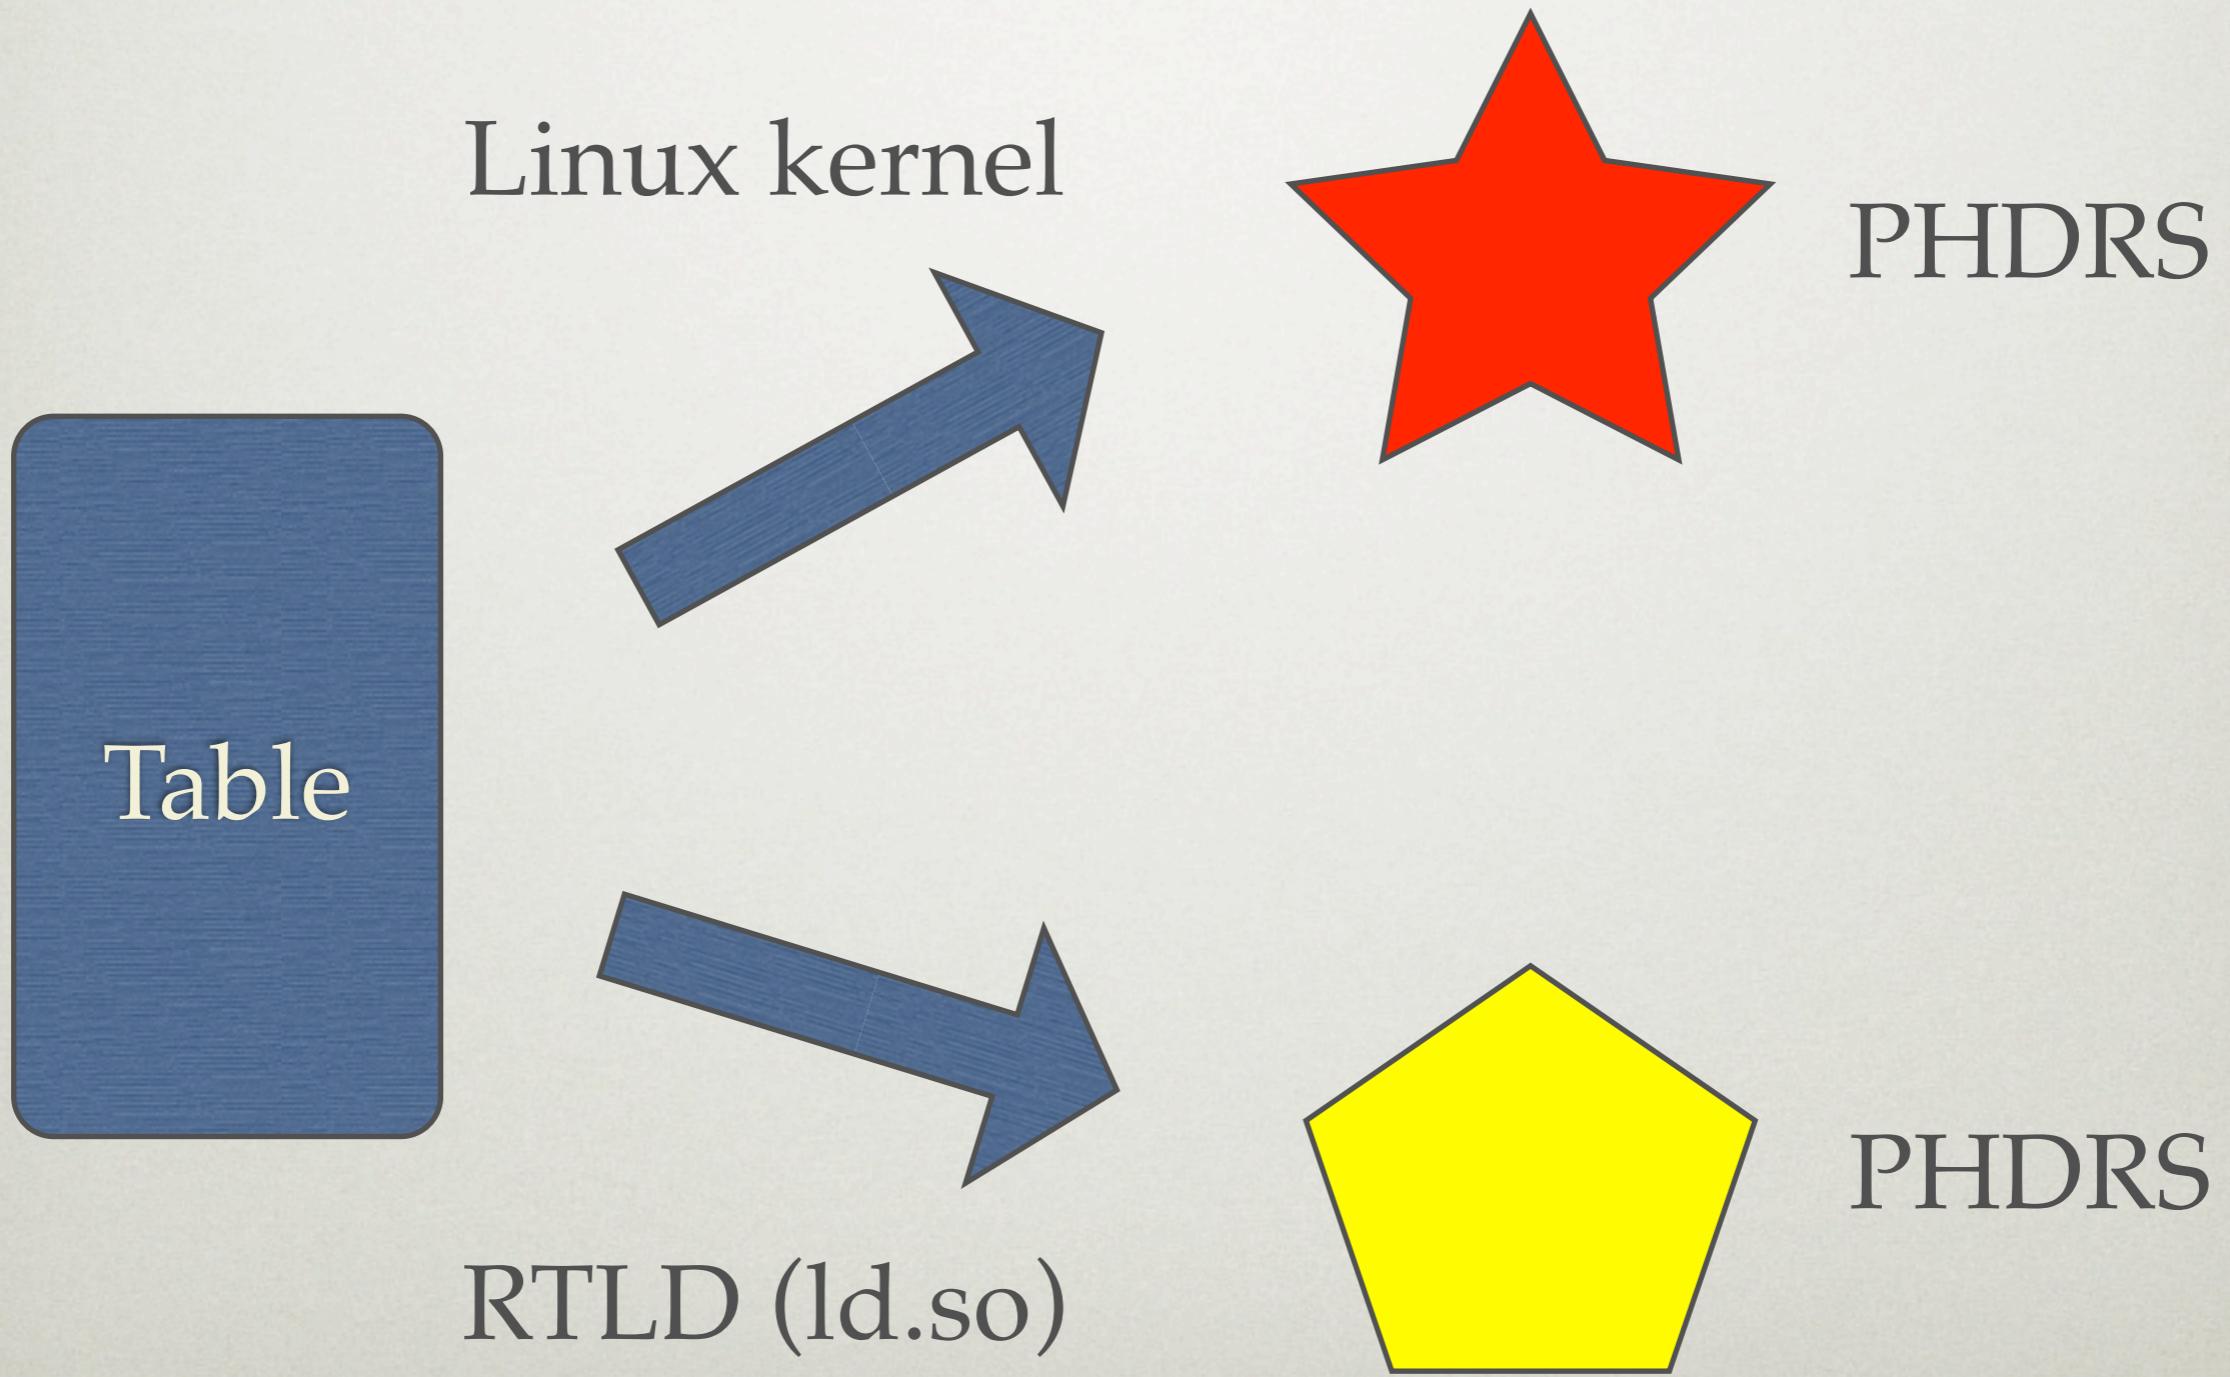
# ...AND NOW WE WAKE UP

---

- One libelf to parse them all?
  - Everyone still builds their own impl.
- “Coming to the B-horror cinema near you!”
  - **Revenge of the Design Committee**
  - **Assault of the Performance Hack**
  - **The Night of the Unchecked Assumption**

# DOUBLE THE PARSING, DOUBLE THE FUN

---



# ELF HEADER

---

- Generic information  
(Version,  
architecture, file  
type, etc.)
- Refers to section  
table & program  
header table

```
typedef struct {
    unsigned char      e_ident[16];
    Elf32_Half        e_type;
    Elf32_Half        e_machine;
    Elf32_Word        e_version;
    Elf32_Addr        e_entry;
    Elf32_Off          e_phoff;
    Elf32_Off          e_shoff;
    Elf32_Word        e_flags;
    Elf32_Half        e_ehsize;
    Elf32_Half        e_phentsize;
    Elf32_Half        e_phnum;
    Elf32_Half        e_shentsize;
    Elf32_Half        e_shnum;
    Elf32_Half        e_shstrndx;
} Elf32_Ehdr;
```

# SECTION HEADERS

---

- Label a range of offsets
- Represent 2 views of the file
  - Offset view
  - Address view
- **Labels for byte subsets** of the file

```
typedef struct
{
    Elf32_Word sh_name;
    Elf32_Word sh_type;
    Elf32_Word sh_flags;
    Elf32_Addr sh_addr;
    Elf32_Off  sh_offset;
    Elf32_Word sh_size;
    Elf32_Word sh_link;
    Elf32_Word sh_info;
    Elf32_Word sh_addralign;
    Elf32_Word sh_entsize;
} Elf32_Shdr;
```

# SECTION HEADER FUN

---

- Do not guarantee non-overlay
  - “What is at offset/address x mean?”
  - Context sensitive
  - We could just reject overlap
  - Yet another obfuscation talk?
- Loader has to parse large section table
  - OMG 30 entries!

# PROGRAM HEADERS (SEGMENTS)

---

- Loader ignores sections
- Has separate ‘program header table’.
- A **different label system**
- != x86 hardware segments

```
typedef struct
{
    Elf32_Word p_type;
    Elf32_Off  p_offset;
    Elf32_Addr p_vaddr;
    Elf32_Addr p_paddr;
    Elf32_Word p_filesz;
    Elf32_Word p_memsz;
    Elf32_Word p_flags;
    Elf32_Word p_align;
} Elf32_Phdr;
```

# IDA USES SECTIONS, KERNEL USES SEGMENTS

```
; ===== SUBROUTINE =====  
  
; Attributes: bp-based frame  
  
main    public main  
        proc near             ; DATA XREF: _start+1D↑  
            push    rbp  
            mov     rbp, rsp  
            mov     edi, offset hello_2 ; "Hello, World\n"  
            call    _puts  
            pop     rbp  
            retn  
        endp  
  
;
```

```
$ ./hello  
dlrow ,olleH
```

# MEMORY VS. FILE

---

- Most parsers **read()** file
- Some work on the memory **mmap()**-ed version
- Assume almost **linear** map
- The file describes its own memory layout





Thursday, September 12, 13

# PROGRAM HEADERS (I)

---

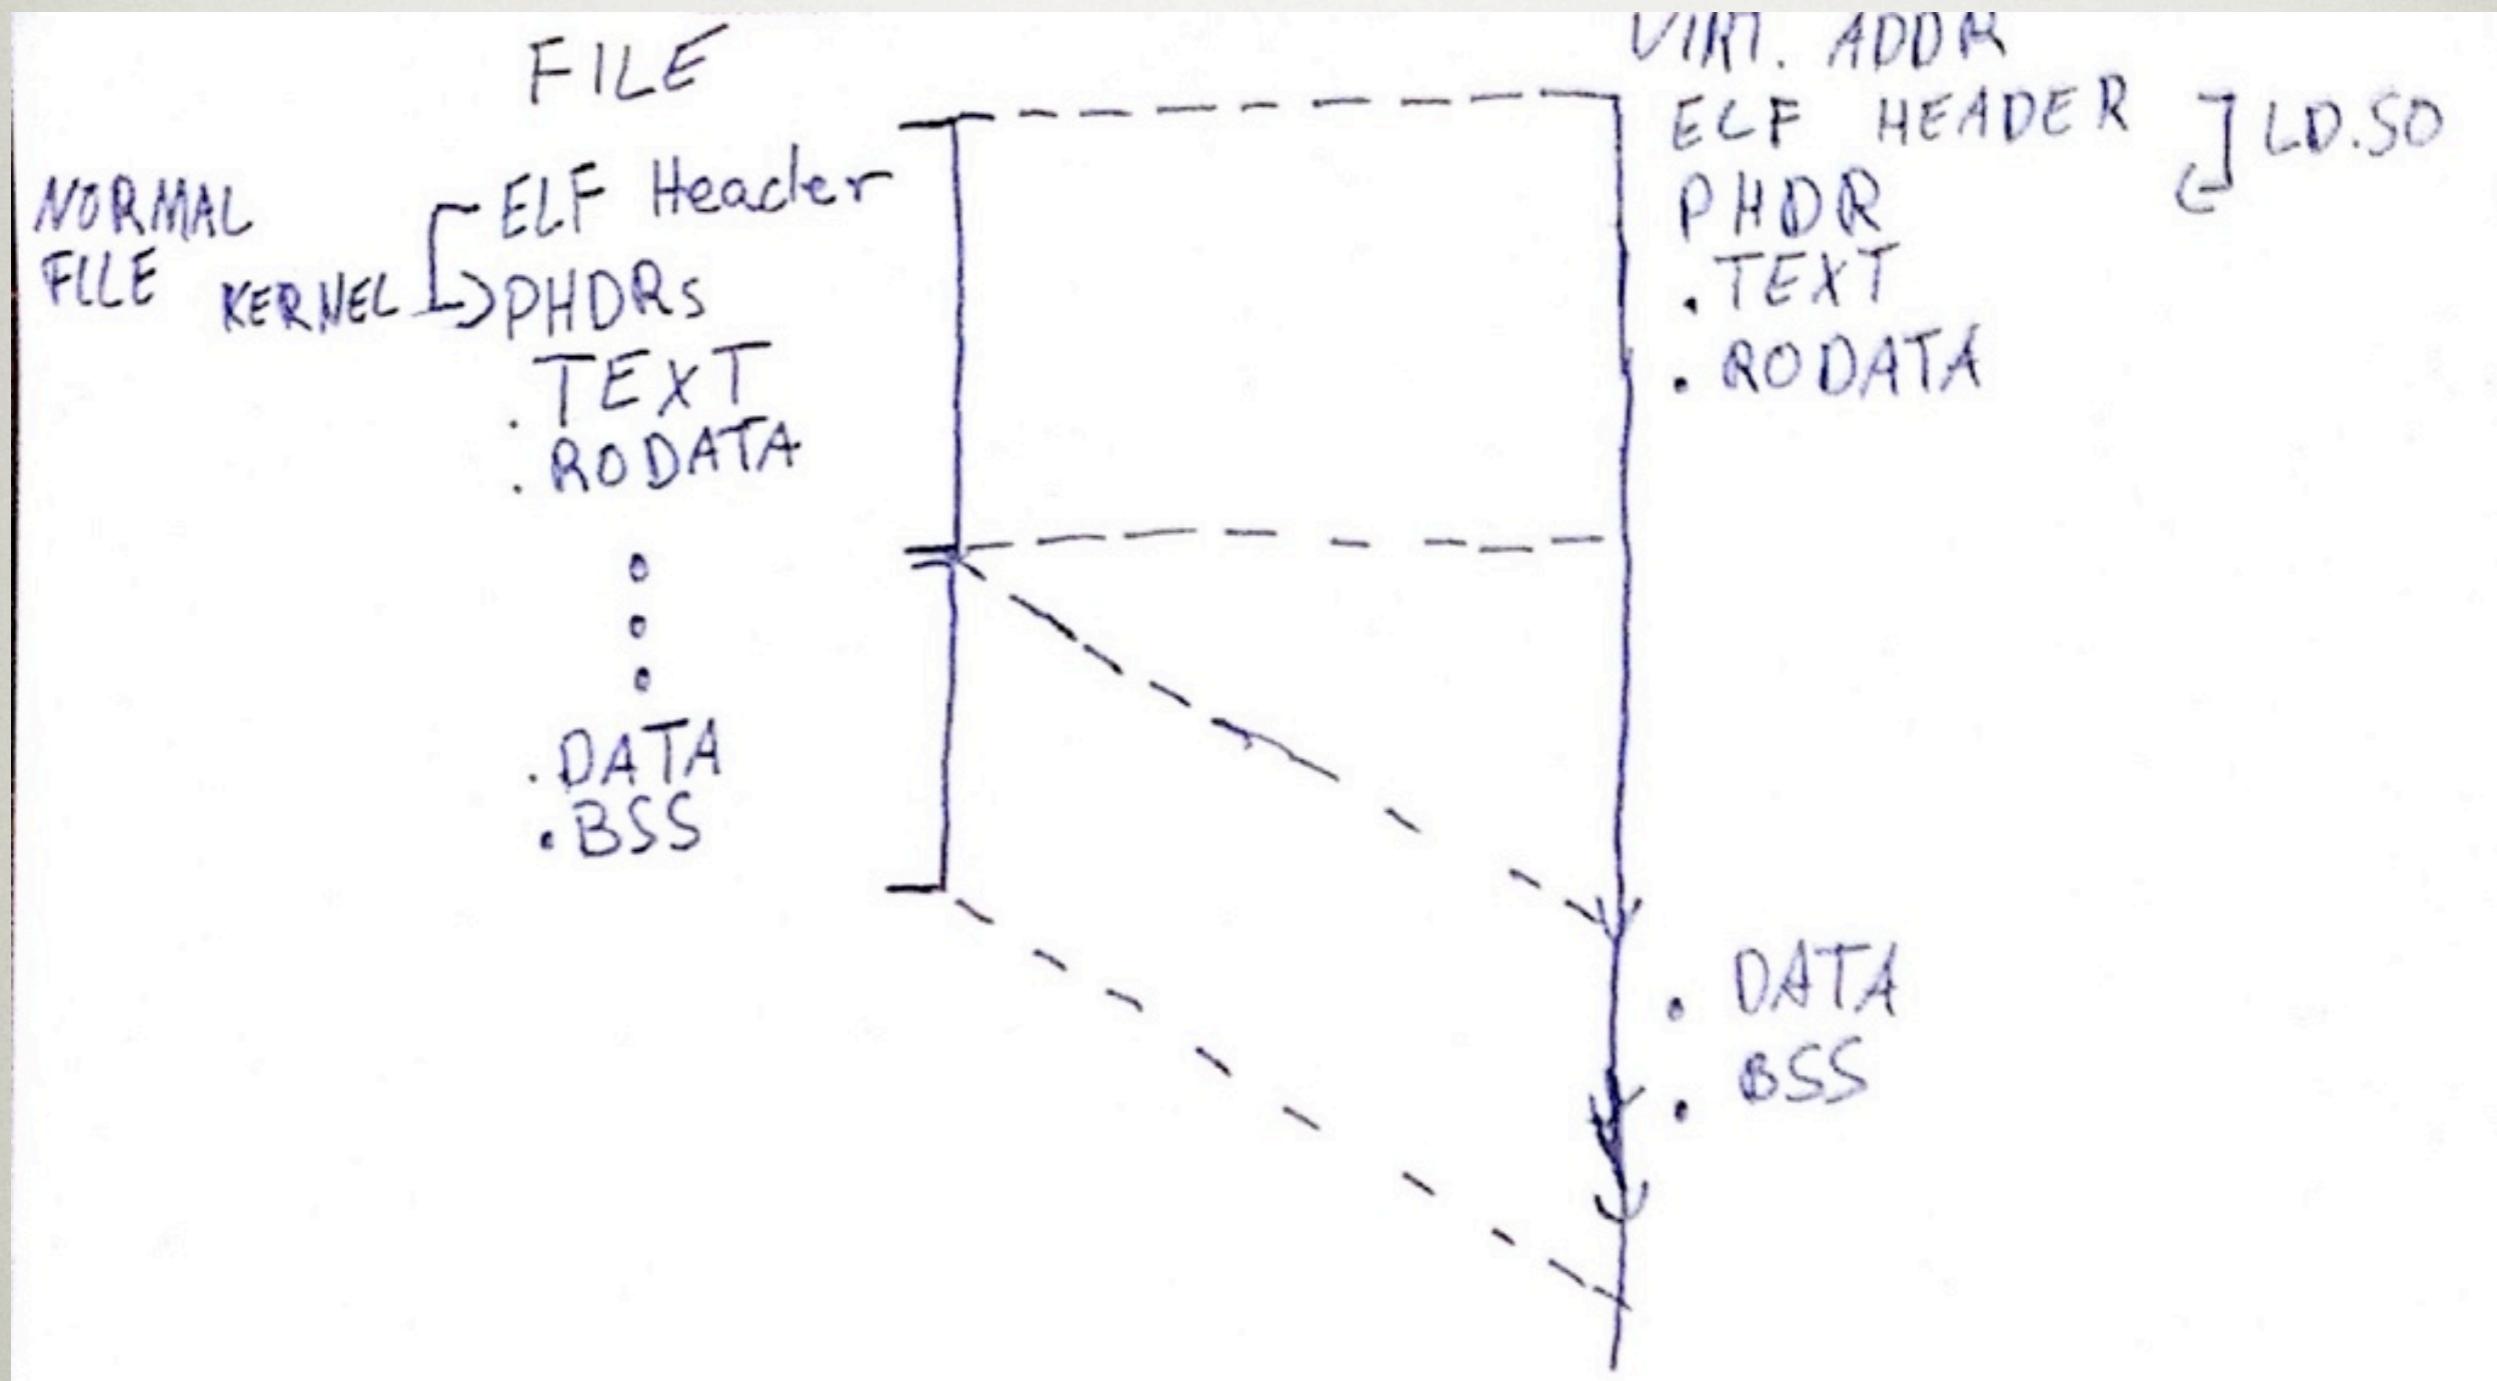
- PT\_LOAD:
- “Map these bytes (+ some /dev/null) to this address range”
- Overlap: each PT\_LOAD is a `mmap()`
- Evad3rs noted that `mmap` replaces existing mappings
- Order of PT\_LOAD is important, but that’s not in the spec.

# THE BIRTH OF AN ELF

---

1. Kernel reads PHDR table into buffer
2. One mmap per PT\_LOAD
3. Loads PT\_INTERP
4. Sends (addr of first PT\_LOAD)+  
hdr.phoff to loader (“AUX vector”)
5. Loader looks for & parses PHDRs  
starting at the address in aux vector

# LD.SO PHDR



# OBLIGATORY KERNEL SOURCE

---

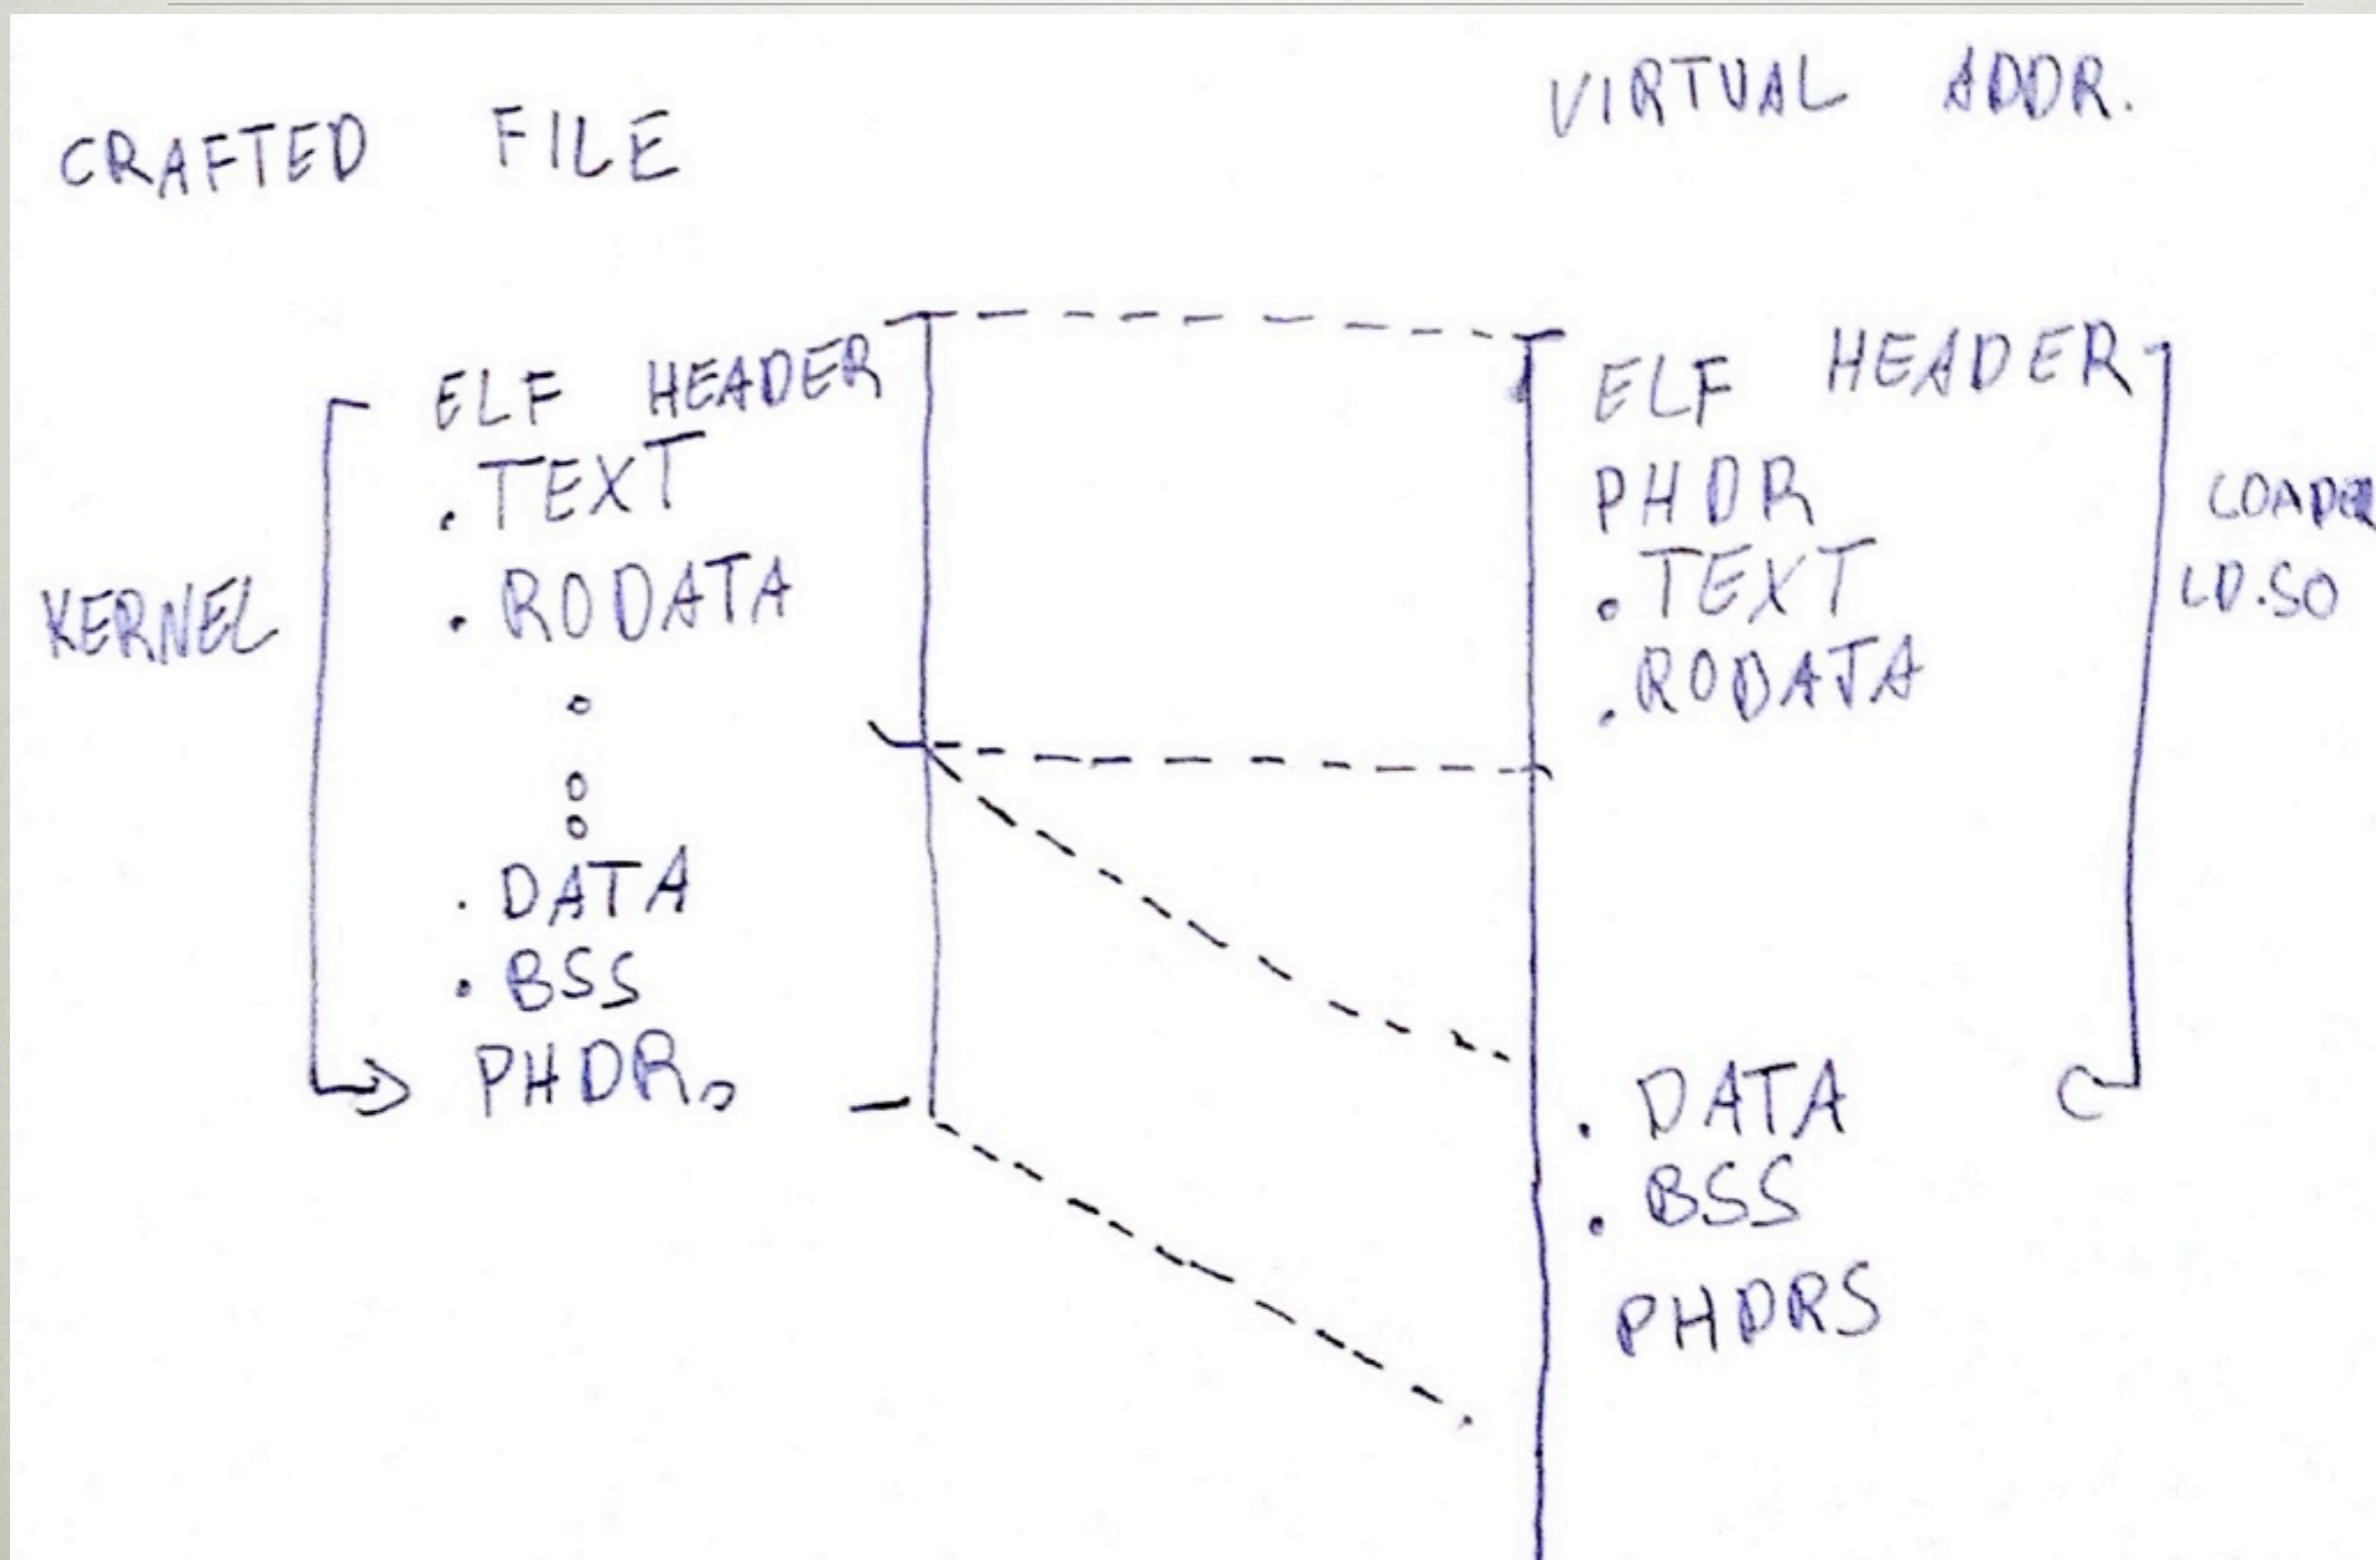
```
/*Redacted from binfmt_elf.c, Linux 3.4, GPLv2*/
elf_phdata = kmalloc(size, GFP_KERNEL);
kernel_read(bprm->file, loc->elf_ex.e_phoff, (char *)
    *elf_phdata, size)
for(phdrs...) {
    /*..., scans for PT_INTERP and PT_LOAD*/
    if(phdr->p_type == PT_LOAD)
        if (!load_addr_set) {
            load_addr = (elf_ppnt->p_vaddr - elf_ppnt-
>p_offset);
            load_addr_set = 1;
        }
    }
NEW_AUX_ENT(AT_PHDR, load_addr + exec->e_phoff);
/*^^^^^ FAIL! */
```

# PHDR BUG

---

- Works Just Fine™ if program headers mapped as first segment.
- Otherwise, points somewhere else
- We can “finger-paint” memory to our liking, and ld.so will use **different PHDRS**
- Hide your PHDRs, the reverse engineers are coming!

# CRAFTED FILE



# WHAT DO WE GET?

---

- Can't map memory, the kernel does that
- PT\_DYNAMIC

# DYNAMIC SECTION

---

- The loader gets symbol tables, relocations, etc., from the `.dynamic` section
  - which it finds through the `PT_DYNAMIC` PHDR
- Only cares about **address**, not file offsets or sections

# HOW TOOLS FIND DYNAMIC

---

- Program Header, by virtual address or offset
- Section
- Symbol

# DYNAMIC ENTRIES

---

- DT\_STRTAB-> Strings (e.g. symbols)
- DT\_SYMTAB -> Symbols
- DT\_RELAT -> Relocations (ASLR, shlib)
- DT\_NEEDED -> imported shared libraries

# SHARED LIBRARIES

---

- Shared libraries (and PI executables) are loaded at a random address
- Allocate space for entire library, then remap
- Size: `End_of_last_PHDR - start_of_first_PHDR`
- **Deranged PT\_LOADs** (ranges “normally” assumed sorted by loader), integer overflow?
- Can change memory permissions elsewhere in program

# “MEMORY FINGER- PAINTING”

---

- PHDRS control vma structs in kernel
- Different paths: mmap(), mmap()->mprotect(), in-kernel parser
- Page-granularity of vmas not reflected in format
- Rounding in various places, overlap
- Combining different paintings

# REWRITING PROGRAMS IN MEMORY WITH LD.SO?

---

- We control **ld.so**'s idea of all relevant sections: GOT, dyn symbols, ...
- **ld.so** resolves (what it thinks are) symbols, writes (what it thinks is) GOT
- Fun way to rewrite your program via crafted .dynamic + library symbols

# SYMBOLS

---

- Symbol is a name for addresses(.so and +x) or offsets(everything else). That is a convention, tools might break if you mix
- **.dynsym** (loader) vs **.syntab** (linker,debugger)
- Undefined symbols, imported from somewhere else
- Weak symbols for ‘default values’
- Can hide symbols from linker with versioning

# RELOCATIONS

---

- Fill in missing references to other objects
- *Theory*: ‘Replace this bit of program / data with some computation on this symbol’
- *Reality*: Read memory(.rela, .dynsym), compute, write memory -- **anywhere**
- @bxsays: **Turing machine!** (see elf-bf-tools on Github, our talk noon tomorrow)

# WHAT IS LINKING?

---

- We are working on formalizing this
- Roughly:
  - Combine sections from different ELF objects
  - Resolve symbols
  - Apply relocations
- Code at the end of the talk

# BUT IT HAS TO BE FAST

---

- Symbol lookups by walking string table are inefficient
- DT\_HASH and DT\_GNU\_HASH
  - hashtable Symbol name -> string

# DT\_HASH

---

- Hashtable, simple hash
- Each hash bucket can point to one symbol
- chain has ‘next’ link for each symbol

uint32_t nbucket
uint32_t nchain
uint32_t bucket[nbucket]
uint32_t chain[nchain]

# WALKING THE HASH TABLE

```
/*dl-lookup.c:268*/  
for (symidx = map->l_buckets[*old_hash %  
map->l_nbuckets];  
     symidx != STN_UNDEF;  
     symidx = map->l_chain[symidx])  
{  
    sym = check_match (&syntab[symidx]);  
    if (sym != NULL)  
        goto found_it;  
}
```

What happens if  
these checks never  
succeed?

# DT\_HASH\_LOOPING: “THE RED QUEEN BREAK-POINT”

---

- Collision entry points to itself
- Resolver hangs
- Without hash(or DT\_GNU\_HASH) it will work fine
- Can make only some symbols hang
- Breakpoints

uint32_t nbucket
uint32_t nchain
uint32_t bucket[nbucket]
uint32_t chain[nchain]

DER TRAGÖDIE ZWEITER TEIL

IN SOVIET RUSSIA,  
MITHRIL FORGES ELF

# INSPIRATION: ELFSH

---

- ERESI reversing framework by Julien Vanegue and team
- .bx and I are big fans!
- ELF handled with libelfsh
- just went back to active development;  
ELF parsing developed as needed =>  
some bugs

# WHAT TO DO DIFFERENTLY

---



- ERESI modifies the binary as little as possible
- In-memory representation close to on-disk
  - Rewrites much of the binary on change
  - Sometimes (inevitably), misses a dependency and breaks the binary

# MITHRIL

---

- Elves smithe powerful swords and chain armour from mithril
- We cut and link ELFes with Mithril
- Parse ELF into a ‘context-free’ representation, write representation into ELF
- Files differ, but semantically identical

# MITHRIL

---

- Ruby prototype
- Rewrites Ubuntu 12.04 64-bit - except
- [github.com/jbangert/mithril](https://github.com/jbangert/mithril)

# BINARYSMITHING 101

---

- Injecting code, ordinary swords:
  - Shift section headers, inject section
  - Shift sections in file, inject data
  - Update PHDRS
- Injecting with Mithril:

```
mysect = Elf::Progbits.new(".inject")
mysect.data.write "Foo\0"
mysect.flags = SHF::SHF_ALLOC
file.progbits.push mysect
```

# GRADUATE BINARYSMITHING

---

- Inject symbol:

```
file.symbols << Symbol.new(...)
```

- Inject relocation:

```
file.relocations << Relocation.new()
```

- Mithril::Writer.to\_file("foo",file)
- Mithril takes care of the rest

# LET'S ROLL OUR OWN LINKER

---

- Showcase the conciseness of Mithril
- What does it really mean to ‘link’ a binary? Upper bound by construction

# USELESS METRICS !

---

- sloccount binutils-2.20:
  - ld:115138, bfd:317996
  - gold:56673 is a better comparison..
- wc -l bin/mithril-ld : 76
- sloccount mithril-0.0.1: 2771

# INVOKING IT

---

```
file = Elf::linkfiles(Elf::newfile,  
ARGV[1..-1].map{|x| Elf::Parser::from_file(x)})  
  
Elf::Writer.to_file(ARGV[0], file)
```

# LET'S HAVE A SOURCE WALKTHROUGH

---

```
require 'mithril'
module Elf
def self.newfile
    ElfFile.new.tap { |x|
        x.filetype = ElfFlags::Type::ET_EXEC
        x.machine = ElfFlags::Machine::EM_X86_64
        x.version = ElfFlags::Version::EV_CURRENT
        x.flags = 0
        x.bits = 64
        x.endian = :little
        x.interp = "/lib64/ld-linux-x86-64.so.2"
    }
end
```

# STATE

---

```
def self.linkfiles(outfile, infilenames)
    progbits = {}
    symbols = {}
    sect_map = {}
    sect_offsets = {}
    symbols = {}
```

```
infiles.each { |infile|
    #link actual data
    (infile.progbits + infile.nobits).each
    { |inbit|
        outbit = (progbits[inbit.name] ||=
Elf::ProgBits.new(inbit.name))
        outbit.flags |= inbit.flags
        outbit.align = [inbit.align,
outbit.align].max
        outbit.sect_type = inbit.sect_type
        sect_offsets[inbit] =
outbit.data.tell
        sect_map[inbit] = outbit
        outbit.data.write inbit.data.read
#TODO: Handle align
    }
}
```

---

```
infile.symbols.each { |symbol|
  next if symbol.name.empty?
  next unless [Elf::STT::STT_OBJECT, Elf::STT::STT_FUNC,
Elf::STT::STT_COMMON, Elf::STT::STT_NOTYPE].include? symbol.type

  if (symbols.include? symbol.name)
    next if symbol.weak? or symbol.undefined?
    raise ArgumentError.new "Duplicate definition of symbol
#{symbol.name}" unless symbols[symbol.name].weak? or
symbols[symbol.name].undefined?
  end
  symbols[symbol.name] = symbol.clone.tap{|outsym|
    unless outsym.undefined?
      outsym.sectoffset += sect_offsets[outsym.section]
      outsym.section = sect_map[outsym.section]
    end
  }
}
```

---

```
progbits.values.each { |proabit| outfile.progbits
<< proabit    }
    symbols.values.each { |sym| outfile.symbols <<
sym}

inf files.map(&:relocations).flatten.each { |rela|
    outfile.relocations << rela.clone.tap{ |x|
        x.offset += sect_offsets[x.section]
        x.section = sect_map[x.section]
        name = x.symbol.name
        x.symbol = symbols[name]
        x.is_dynamic = true
        raise ArgumentError.new "Undefined symbol
#{name}" if(x.symbol.nil?)
    }
}
```

---

```
raise "Undefined entry point  
      _start" unless symbols.include?  
      "_start"  
      entry = symbols["_start"]  
      entry.section.addr = 0x40000  
      outfile.entry =  
      entry.sectoffset +  
      entry.section.addr  
      return outfile
```

# MISSING FEATURES

---

- Section Symbols are handled poorly
  - (That currently breaks crt1.o, which has the `_start` for libc)
- No STF\_MERGE, STT\_COMMON, other ‘unusual’ extensions
  - Sorry, we don’t link Fortran here

# WHAT TO DO WITH IT

---

- Normal ld works, so keep using it
- But load up mithril in Pry / IRB and play around with binaries
- Metadata tricks don't have to be arcane!

# SHOUTOUTS!

---

- LANGSEC conspiracy (<http://langsec.org/>)
- Corkami <http://code.google.com/p/corkami/>
- ERESI project team, past and present
- TinyELF <http://www.muppetlabs.com/~breadbox/software/tiny/teensy.html>
- The Grugq & Silvio Cesare  
(for hooking Sergey on ELF)
- Phrack \cite{Phrack XX:XX}

# QUESTIONS?

---

- @julianbangert, @sergeybratus
- [github.com/jbangert/mithril](https://github.com/jbangert/mithril)
- gem ‘elf-mithril’
- Needs gem ‘bindata’, [github: ‘jbangert / bindata’](https://github.com/jbangert/bindata) (Fork not merged back yet)
- See also USENIX WOOT 2011, 2013 for our “weird machines” papers